

1 Voice processing — calls into records

Speech recognition	Open model + 20 h of 119 call training
Trained portion	Only 3.6% of the model (LoRA)
Extraction	Korean encoder + custom layers
Output	15 fields · always the same format

15 judges · E-Gen severe-disease groups — reperfusion, brain hemorrhage surgery, aortic emergency, severe burns + 11 more

Where it's used

- 119 EMS — the run report is written automatically from the call
- ER — pre-arrival calls go straight to the triage screen
- Regional dispatch — share once instead of re-explaining to 10 hospitals

2 Reliability model — the shelf life of information


Lifespan learning never reads the content. It learns only **who · how often · how** a source has edited its data, and outputs the probability that it is stale right now.

Prior theory — “What changed often has probably changed again”
→ rank scores only, one rule for every source

GoldenLink — “Given this history, how likely is it stale now?”

On the same history, change-frequency methods show up to **168x** higher probability error (IBS)

Where it's used · validated in 4 domains



The 0.48 floor for maps · places is near random (0.5) — such data is screened out first by the **automatic diagnosis** (applicability check) before adoption.

AI search · RAG is validated by outcome, not discrimination — stale-evidence answers 33.3% → 23.3% (preliminary, single corpus)

Discrimination: 1.0 = perfect · 0.5 = random. Paper “Learning the Shelf Life of Information” submitted to an international journal.

We put an expiry date on data

The engine proven on ER beds, applied to any data that goes stale

Today, all data is refreshed on the same cycle and trusted equally until the next refresh. But every dataset goes stale at its own speed — even within one feed, item lifespans differ **8x**.

Information Reliability Engine

API · On-premise

IN	You send · data IDs and edit timestamps — no content or personal data
OUT	You get · probability it's correct now · expected lifespan · when to re-check
START	Before adoption · we start with an automatic diagnosis of whether it works on your data

What it does

- Tells you when to re-check** — 11% less time on stale data, same budget
- Filters before use** — low probability means verify first
- Audits reported data** — reports where agency-reported values go stale

Who it's for

Fire · dispatch · local gov't	Bed reliability + automatic call records (GoldenLink) — choose hospitals from trustworthy candidates
AI search · chatbot teams	Checks whether evidence is still valid before answering, re-verifying only the sources that need it — stale-evidence answers 33.3% → 23.3%
Data providers · platforms	Per-item re-check alerts + reports on where reported values go stale — 11% less time on stale data, same budget (median)

Team Junhyeok Choi Product Engineer · Seungju Lee Security · Taewoo Kim Forward Deployed Engineer · Donghyun Kim Data · Hoyoung Kwak ML



INSTAGRAM @playground
EMAIL edclplayground@gmail.com

INSTAGRAM

Delayed ER transfers,
a.k.a. ‘ER ping-pong’

Why do ambulances wander in search of a hospital?

2.4x

Two-year rise in transfers taking 30+ minutes from scene arrival to hospital departure



Is it really a lack of beds?

Open to find out. →

TEAM PLAYGROUND

AI FESTA 2026